

CMPE58C: Sp. Tp. Mobile Location Tracking & Motion Sensing

Data Fusion

Can Tunca

-


Fall 2022

Lecture Overview

- Why “Data Fusion” and not “Sensor Fusion”
 - Sensor fusion implies fusion of raw sensor data
 - Fusion can take place at higher levels, e.g. fusing the output of two positioning methods
- Some topics that will be covered:
 - Bayesian filters introduction
 - Kalman filter
 - Linearized Kalman filters
 - Extended Kalman filter
 - Unscented Kalman filter
 - Particle filters
 - Data fusion examples
 - 1D positioning (velocity + absolute measurements)
 - 2D positioning (odometry)
 - Pedestrian dead reckoning

Data Fusion: Reasons

- Handling uncertainty in raw measurements (noise/bias)
- Desire to fuse measurements from multiple sensors
- Incorporating accuracy estimates of measurements (if exists)
- Smoothing (preventing fluctuations)
- Modeling error and noise of different sources




Bayesian Filters

- A probabilistic framework for recursive state estimation
- Estimating the probability distribution over the space of all possible states at time t

$$\underset{\substack{\nearrow \\ \text{belief}}}{\text{Bel}(\mathbf{x}_t)} = p(\mathbf{x}_t | \underbrace{\mathbf{z}_1, \dots, \mathbf{z}_t}_{\substack{\nwarrow \\ \text{measurements made up to } t}})$$

- Representing state as a probability distribution allows us to provide both an estimate and its uncertainty
- i.e. the belief can be queried to get the probability that \mathbf{x}_t is the true state
- The complexity of belief increases as measurements accumulate, to combat this Bayesian filters assume the underlying dynamic system is a **Markov process**

Markov Process



- Two assumptions
 - The current state \mathbf{x}_t depends on only the previous state $\mathbf{x}_{t-\delta t}$.
 - A measurement \mathbf{z}_t depends on only the current state \mathbf{x}_t
- In other words, all the information required to estimate state at time t is given by the information available at time $t - 1$: aka Markovian property
- Such systems are defined by two distributions:

$$p(\mathbf{x}_t | \mathbf{x}_{t-\delta t})$$

propagation
model

$$p(\mathbf{z}_t | \mathbf{x}_t)$$

measurement
model

Belief Propagation and Correction

- Under the Markov assumption current state updated based on previous state:

$$\text{Bel}^-(\mathbf{x}_t) = \int p(\mathbf{x}_t | \mathbf{x}_{t-\delta t}) \text{Bel}(\mathbf{x}_{t-\delta t}) d\mathbf{x}_{t-\delta t}$$

- This is called the estimated belief or *prior* distribution
- The prior is then is corrected to obtain *posterior* distribution:

$$\text{Bel}(\mathbf{x}_t) = \alpha_t p(\mathbf{z}_t | \mathbf{x}_t) \text{Bel}^-(\mathbf{x}_t) \quad (\alpha_t \text{ is a normalization factor to make sure Bel is a prob. distribution})$$

- Different measurement models can be used to incorporate different measurement types (Bayesian filters naturally support sensor fusion)
- This is a probabilistic “framework”, the concrete implementation should choose how to represent the belief. We’ll see the most common implementations next...

Kalman Filter

- Assuming belief is Gaussian distributed for all time t , hence representable by a mean and covariance
- State transition (propagation) model is linear:

$$\mathbf{x}_t = \mathbf{F}_t \mathbf{x}_{t-\delta t} + \mathbf{w}_t$$

state transition matrix zero mean Gaussian noise
(process noise)


- Measurements are related to the state by a linear function:

$$\mathbf{z}_t = \mathbf{H}_t \mathbf{x}_t + \mathbf{v}_t$$

matrix relating measurement to state zero mean Gaussian noise (measurement noise)

Kalman Filter

- How to compute belief mean and covariance?
- Two stages: **PREDICT** and **CORRECT**, a recursive process
- Correspond to state transition and measurement models, respectively



PREDICT

$$\hat{\mathbf{x}}_{t|t-1} = \mathbf{F}_t \hat{\mathbf{x}}_{t-1|t-1} + \mathbf{B}_t \mathbf{u}_t$$

$$\mathbf{P}_{t|t-1} = \mathbf{F}_t \mathbf{P}_{t-1|t-1} \mathbf{F}_t^T + \mathbf{Q}_t$$

CORRECT

$$\hat{\mathbf{x}}_{t|t} = \hat{\mathbf{x}}_{t|t-1} + \mathbf{K}_t (\mathbf{y}_t - \mathbf{H}_t \hat{\mathbf{x}}_{t|t-1})$$

$$\mathbf{K}_t = \mathbf{P}_{t|t-1} \mathbf{H}_t^T (\mathbf{H}_t \mathbf{P}_{t|t-1} \mathbf{H}_t^T + \mathbf{R}_t)^{-1}$$

$$\mathbf{P}_{t|t} = (\mathbf{I} - \mathbf{K}_t \mathbf{H}_t) \mathbf{P}_{t|t-1}$$

$\hat{\mathbf{x}}$: Estimated state.

\mathbf{F} : State transition matrix (i.e., transition between states).

\mathbf{u} : Control variables.

\mathbf{B} : Control matrix (i.e., mapping control to state variables).

\mathbf{P} : State variance matrix (i.e., error of estimation).

\mathbf{Q} : Process variance matrix (i.e., error due to process).

\mathbf{y} : Measurement variables.

\mathbf{H} : Measurement matrix (i.e., mapping measurements onto state).


\mathbf{K} : Kalman gain.

\mathbf{R} : Measurement variance matrix (i.e., error from measurements).

Kalman Filter: Predict

1 $\hat{\mathbf{x}}_{t|t-1} = \mathbf{F}_t \hat{\mathbf{x}}_{t-1|t-1} + \mathbf{B}_t \mathbf{u}_t$

2 $\mathbf{P}_{t|t-1} = \mathbf{F}_t \mathbf{P}_{t-1|t-1} \mathbf{F}_t^T + \mathbf{Q}_t$




- Predict stage advances time ($t-1 \rightarrow t$)
- State transition to get an “a priori” estimate for the next timestep
- **1** computes the a priori mean (the actual estimate)
- **2** computes the a priori error covariance (the error of the estimate)
- Together they form a Gaussian representing the belief
- \mathbf{u}_t is the optional control input, some external known factor affecting state
 - Either a controllable action (e.g. throttle in a car) or a measurable quantity advancing the state (e.g. acceleration), usually denoting a known **change** to the system - A possible place for data fusion!
 - \mathbf{Q}_t embeds both error due to state transition and control input

Kalman Filter: Correct

1 $\hat{\mathbf{x}}_{t|t} = \hat{\mathbf{x}}_{t|t-1} + \mathbf{K}_t (\mathbf{y}_t - \mathbf{H}_t \hat{\mathbf{x}}_{t|t-1})$


2 $\mathbf{K}_t = \mathbf{P}_{t|t-1} \mathbf{H}_t^T (\mathbf{H}_t \mathbf{P}_{t|t-1} \mathbf{H}_t^T + \mathbf{R}_t)^{-1}$

3 $\mathbf{P}_{t|t} = (\mathbf{I} - \mathbf{K}_t \mathbf{H}_t) \mathbf{P}_{t|t-1}$



- Corrects estimate at time t with a measurement
- Computes the “posterior” estimate for current time
- 1 computes the posterior mean (the actual estimate)
- 3 computes the posterior error covariance (the error of the estimate)
- 2 determines how to fuse the prior belief and measurement: **Kalman Gain**
 - We shouldn't trust new measurement fully, otherwise we wouldn't be using accumulated knowledge
 - We shouldn't trust the previously predicted state, otherwise we won't be correcting by new measurements
 - Quite like the blending weight in complementary filter
 - Computed optimally, error models \mathbf{Q}_t and \mathbf{R}_t influence the outcome

Kalman Filter Example: 1D Positioning



$$y_t = y_{t-1} + v_{t-1}\Delta t$$

$$v_t = v_{t-1}$$


- Assume
 - A car travelling on a straight line
 - Intermittent GPS readings
 - No other measurements (no velocity, no accelerometer)
 - No fusion (only filtering)
 - Our best velocity estimate is the previous one (since no measurement)
- Can we get a smooth estimate of position?
 - Also continue tracking in case of no GPS (e.g. car going in a tunnel)

WE NEED TO DEFINE
THESE MODELS

$$\mathbf{x}_t = \mathbf{F}_t \mathbf{x}_{t-\delta t} + \mathbf{w}_t$$

$$\mathbf{z}_t = \mathbf{H}_t \mathbf{x}_t + \mathbf{v}_t$$

Kalman Filter Example: 1D Positioning



$$y_t = y_{t-1} + v_{t-1}\Delta t$$

$$v_t = v_{t-1}$$

WE NEED TO DEFINE
THESE MODELS


$$\mathbf{x}_t = \mathbf{F}_t \mathbf{x}_{t-\delta t} + \mathbf{w}_t$$

$$\mathbf{z}_t = \mathbf{H}_t \mathbf{x}_t + \mathbf{v}_t$$

$$\mathbf{x}_t = \begin{pmatrix} y_t \\ v_t \end{pmatrix} \begin{matrix} \nearrow \text{position} \\ \nearrow \text{velocity} \end{matrix} \longrightarrow \mathbf{F}_t = \begin{pmatrix} 1 & \Delta t \\ 0 & 1 \end{pmatrix}$$

$$\mathbf{z}_t = \hat{y}_t \begin{matrix} \nearrow \text{position measurement (GPS)} \end{matrix} \longrightarrow \mathbf{H}_t = \begin{pmatrix} 1 & 0 \end{pmatrix}$$

Kalman Filter Example: 1D Positioning



$$\mathbf{F}_t = \begin{pmatrix} 1 & \Delta t \\ 0 & 1 \end{pmatrix} \quad \mathbf{H}_t = \begin{pmatrix} 1 & 0 \end{pmatrix}$$

\mathbf{x}_0 : Set to your initial estimate (could be zeroes): (2x1 matrix)

\mathbf{P}_0 : Set to your initial uncertainty (could be big if unsure) (2x2 matrix)

\mathbf{Q}_t : How much do you trust your state transition model? (2x2 matrix)

\mathbf{R}_t : How much do you trust the measurements (GPS)? (1x1 matrix)

$$y_t = y_{t-1} + v_{t-1} \Delta t$$

$$v_t = v_{t-1}$$

PREDICT

No control input in this example

$$\hat{\mathbf{x}}_{t|t-1} = \mathbf{F}_t \hat{\mathbf{x}}_{t-1|t-1} + \mathbf{B}_t \mathbf{u}_t$$

$$\mathbf{P}_{t|t-1} = \mathbf{F}_t \mathbf{P}_{t-1|t-1} \mathbf{F}_t^T + \mathbf{Q}_t$$


CORRECT

$$\hat{\mathbf{x}}_{t|t} = \hat{\mathbf{x}}_{t|t-1} + \mathbf{K}_t (\mathbf{y}_t - \mathbf{H}_t \hat{\mathbf{x}}_{t|t-1})$$

$$\mathbf{K}_t = \mathbf{P}_{t|t-1} \mathbf{H}_t^T (\mathbf{H}_t \mathbf{P}_{t|t-1} \mathbf{H}_t^T + \mathbf{R}_t)^{-1}$$

$$\mathbf{P}_{t|t} = (\mathbf{I} - \mathbf{K}_t \mathbf{H}_t) \mathbf{P}_{t|t-1}$$

Kalman Filter Example: 1D Positioning with Acceleration



$$y_t = y_{t-1} + v_{t-1}\Delta t + \frac{1}{2}a_t\Delta t^2$$
$$v_t = v_{t-1} + a_t\Delta t$$


- Assume
 - We can also measure acceleration
 - Same state elements (position + velocity)
- Can we use it to improve our estimate?
- How to incorporate it into our model?
 - Control input is a good place

WE NEED TO DEFINE
THESE MODELS

$$\mathbf{x}_t = \mathbf{F}_t \mathbf{x}_{t-\delta t} + \mathbf{w}_t$$

$$\mathbf{z}_t = \mathbf{H}_t \mathbf{x}_t + \mathbf{v}_t$$

Kalman Filter Example: 1D Positioning with Acceleration



$$B_t = \begin{bmatrix} \Delta t^2/2 \\ \Delta t \end{bmatrix} \quad u_t = a_t$$

We should also adapt Q_t to incorporate acceleration error

Everything else is the same

This is data fusion! (we've fused accelerometer and GPS)

$$y_t = y_{t-1} + v_{t-1}\Delta t + \frac{1}{2}a_t\Delta t^2$$
$$v_t = v_{t-1} + a_t\Delta t$$

PREDICT

$$\hat{\mathbf{x}}_{t|t-1} = \mathbf{F}_t \hat{\mathbf{x}}_{t-1|t-1} + \mathbf{B}_t \mathbf{u}_t$$

$$\mathbf{P}_{t|t-1} = \mathbf{F}_t \mathbf{P}_{t-1|t-1} \mathbf{F}_t^T + \mathbf{Q}_t$$

CORRECT


$$\hat{\mathbf{x}}_{t|t} = \hat{\mathbf{x}}_{t|t-1} + \mathbf{K}_t (\mathbf{y}_t - \mathbf{H}_t \hat{\mathbf{x}}_{t|t-1})$$

$$\mathbf{K}_t = \mathbf{P}_{t|t-1} \mathbf{H}_t^T (\mathbf{H}_t \mathbf{P}_{t|t-1} \mathbf{H}_t^T + \mathbf{R}_t)^{-1}$$


$$\mathbf{P}_{t|t} = (\mathbf{I} - \mathbf{K}_t \mathbf{H}_t) \mathbf{P}_{t|t-1}$$

Kalman Filter Examples: Some Observations

- Even if I get no GPS for a while, the velocity state element will advance the state
- Control input can be used to incorporate external measurements (first possibility for sensor fusion)
- Different measurement models may be used for incorporating different data (second possibility of sensor fusion)
 - Example: What if we have **speed** measurements? $\mathbf{H}_t = \begin{pmatrix} 0 & 1 \end{pmatrix}$
 - This will also correct position! (correlated state elements)
 - Can specify a different error model \mathbf{R}_t
 - Can be fed into the filter in an alternating fashion (or whenever data is available)
- In a way Kalman filter fuses two Gaussians optimally
 - Prediction error grows until a measurement comes in



Kalman Filter Example: 2D Positioning



$$x_t = x_{t-1} + v_{t-1} \Delta t \cos(\theta_{t-1} + \Delta\theta_t)$$

$$y_t = y_{t-1} + v_{t-1} \Delta t \sin(\theta_{t-1} + \Delta\theta_t)$$

$$\theta_t = \theta_{t-1} + \Delta\theta_t$$

- Assume
 - We can measure speed and heading change (e.g. car speedometer + steering wheel input)
 - Intermittent GPS readings
- We'll revisit this...

WE NEED TO DEFINE THESE MODELS, BUT **CAN WE?**

$$\mathbf{x}_t = \mathbf{F}_t \mathbf{x}_{t-\delta t} + \mathbf{w}_t$$

$$\mathbf{z}_t = \mathbf{H}_t \mathbf{x}_t + \mathbf{v}_t$$

Kalman Filter Conclusion

- Assuming the assumptions hold (linear state transition and measurement models, Gaussian belief and noise), Kalman filter is an optimal filter
- i.e. we can compute the Bayesian filter framework equations exactly
- It is important to set noise covariances right
- Preferred solution for linear or near-linear systems where the belief is unimodal
- Computationally efficient (only a few matrix operations)
- A versatile tool that you can use in many areas (not just for positioning!)

Extended Kalman Filter (EKF)

- What if we can linearize state transition and/or measurement models locally?
- EKF: A common variant of linearized Kalman filters
- It is no longer an optimal filter, an approximation
- The approximation is still representable by a Gaussian
- Works well if the approximation is good enough (i.e. non-linearity is not too much)
- We can then use arbitrary functions for state transition and measurement models, as long as they are differentiable:

$$\mathbf{x}_t = f(\mathbf{x}_{t-1}, \mathbf{u}_t) + \mathbf{w}_t$$

$$\mathbf{z}_t = h(\mathbf{x}_t) + \mathbf{v}_t$$

Extended Kalman Filter

PREDICT

$$\hat{\mathbf{x}}_{t|t-1} = \underline{f(\hat{\mathbf{x}}_{t-1|t-1}, \mathbf{u}_t)} + \mathbf{v}_t$$

$$\mathbf{P}_{t|t-1} = \mathbf{F}_t \mathbf{P}_{t-1|t-1} \mathbf{F}_t^T + \mathbf{Q}_t$$

CORRECT

$$\hat{\mathbf{x}}_{t|t} = \hat{\mathbf{x}}_{t|t-1} + \mathbf{K}_t (\mathbf{y}_t - \underline{h(\hat{\mathbf{x}}_{t|t-1})})$$

$$\mathbf{K}_t = \mathbf{P}_{t|t-1} \mathbf{H}_t^T (\mathbf{H}_t \mathbf{P}_{t|t-1} \mathbf{H}_t^T + \mathbf{R}_t)^{-1}$$


$$\mathbf{P}_{t|t} = (\mathbf{I} - \mathbf{K}_t \mathbf{H}_t) \mathbf{P}_{t|t-1}$$

- The only differences are the models marked by red, rest are the same as KF
- \mathbf{F}_t and \mathbf{H}_t are Jacobian matrices, i.e. first order partial derivatives of \mathbf{f} and \mathbf{h} (see next slide)


Jacobian Matrix

- The matrix of first-order derivatives of a vector function
- First order Taylor expansion
- The orientation of the plane tangent at a given point

$$\mathbf{f} = \begin{pmatrix} f_1 \\ \vdots \\ f_m \end{pmatrix} \quad \mathbf{F} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix}$$



EKF Example: 2D Positioning Revisited



$$x_t = x_{t-1} + v_{t-1} \Delta t \cos(\theta_{t-1} + \Delta\theta_t)$$

$$y_t = y_{t-1} + v_{t-1} \Delta t \sin(\theta_{t-1} + \Delta\theta_t)$$


$$\theta_t = \theta_{t-1} + \Delta\theta_t$$

- Assume
 - We can measure speed and heading change (e.g. car speedometer + steering wheel input)
 - Intermittent GPS readings
- This is a non-linear state transition model, but measurement model is still linear!
 - So we'll need to only compute the Jacobian \mathbf{F}_t

$$\mathbf{x}_t = \begin{pmatrix} x_t \\ y_t \\ \theta_t \end{pmatrix}$$

$$f(\mathbf{x}_{t-1}, \mathbf{u}_t) = \begin{pmatrix} x_{t-1} + v_{t-1} \Delta t \cos(\theta_{t-1} + \Delta\theta_t) \\ y_{t-1} + v_{t-1} \Delta t \sin(\theta_{t-1} + \Delta\theta_t) \\ \theta_{t-1} + \Delta\theta_t \end{pmatrix}$$

EKF Example: 2D Positioning Revisited



$$x_t = x_{t-1} + v_{t-1} \Delta t \cos(\theta_{t-1} + \Delta\theta_t)$$

$$y_t = y_{t-1} + v_{t-1} \Delta t \sin(\theta_{t-1} + \Delta\theta_t)$$

$$\theta_t = \theta_{t-1} + \Delta\theta_t$$

$$\mathbf{x}_{t-1} = \begin{pmatrix} x_{t-1} \\ y_{t-1} \\ \theta_{t-1} \end{pmatrix} \quad f(\mathbf{x}_{t-1}, \mathbf{u}_t) = \begin{pmatrix} x_{t-1} + v_{t-1} \Delta t \cos(\theta_{t-1} + \Delta\theta_t) \\ y_{t-1} + v_{t-1} \Delta t \sin(\theta_{t-1} + \Delta\theta_t) \\ \theta_{t-1} + \Delta\theta_t \end{pmatrix}$$


$$\mathbf{F}_t = \frac{\partial f(\mathbf{x}_{t-1}, \mathbf{u}_t)}{\partial \mathbf{x}_{t-1}} = \begin{pmatrix} 1 & 0 & -v_{t-1} \Delta t \sin(\theta_{t-1} + \Delta\theta_t) \\ 0 & 1 & v_{t-1} \Delta t \cos(\theta_{t-1} + \Delta\theta_t) \\ 0 & 0 & 1 \end{pmatrix}$$

Extended Kalman Filter Conclusion

- Very powerful - allows us to use Kalman filters for many problems!
- But still an approximation...
- Best suited for functions that are near-linear in the short term
- Not suitable for multi-model belief distributions (non-linearity usually leads to such distributions)
- Initial state is more important
 - Re: the previous example: What if we do not know the initial heading of the car?
 - EKF may have trouble estimating it (too much non-linearity)
- Can we do a better approximation?


Unscented Kalman Filter (UKF)

- Another linearized Kalman filter, with a different approximation method
- Rather than linearizing around a single point, use **sigma points**
- Transform each sigma point through the non-linear function (aka Unscented Transform)
- Recompute Gaussian from the transformed and weighted points
- The output is similarly a Gaussian, but potentially a better one




Unscented Transform Example

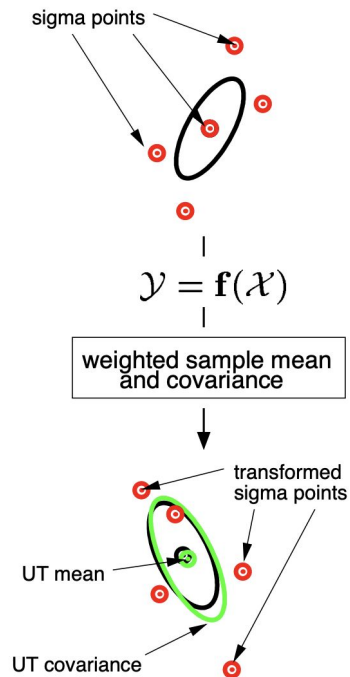
Actual (sampling)



Linearized (EKF)




UT



UKF Considerations


- Sigma points are deterministically selected to represent the Gaussian
- \mathcal{X} is the matrix defining sigma points (i 's are its columns)
- λ influences how far the sigma points are from the mean
- Points are weighted according to their distance from the mean
- UKF deals with non-linearity better, but it's not foolproof
- We still assume Gaussians and unimodal belief distribution

$$\begin{aligned}\mathcal{X}^{[0]} &= \mu \\ \mathcal{X}^{[i]} &= \mu + \left(\sqrt{(n + \lambda) \Sigma} \right)_i \quad \text{for } i = 1, \dots, n \\ \mathcal{X}^{[i]} &= \mu - \left(\sqrt{(n + \lambda) \Sigma} \right)_{i-n} \quad \text{for } i = n + 1, \dots, 2n\end{aligned}$$




Multimodal Belief Example

Imagine a robot which can detect doors, but it doesn't know which door it is




Multimodal Belief Example




After some movement: Note how the estimates widen (incorporating potential error due to motion)

Multimodal Belief Example




After second door detection: Estimate gets sharper
(the peak at the first door is low due to our prior belief)

Multimodal Belief Example



We have now converged on a state we can be confident in

How to represent multimodal belief?



- Such a multimodal distribution cannot be represented accurately by a single Gaussian
- Maybe mixture of Gaussians, but what if measurement model is not Gaussian as well?
(or a complex motion model may disrupt the individual Gaussians)
- We need a generalizable solution for arbitrary distributions
- Next: Particle filter!

Particle Filter


- aka Sequential Monte Carlo
- Represent belief by a set of samples also called “particles”

$$\mathcal{X}_t := x_t^{[1]}, x_t^{[2]}, \dots, x_t^{[M]} \quad \text{M is the number of particles}$$

- Particles are individual “hypotheses”, the likelihood of including a particle in the set should ideally be proportional to the posterior belief


$$x_t^{[m]} \sim p(x_t \mid z_{1:t}, u_{1:t})$$

- Implication: Denser regions should be populated by more particles
- As M goes to infinity we represent the belief better (practically we don't need as much)



How to sample from Belief?

- Belief is an arbitrary distribution, could be multimodal and complex
- So, sampling from such a distribution is not a straightforward task





How do we generate samples from f ?

Importance Sampling

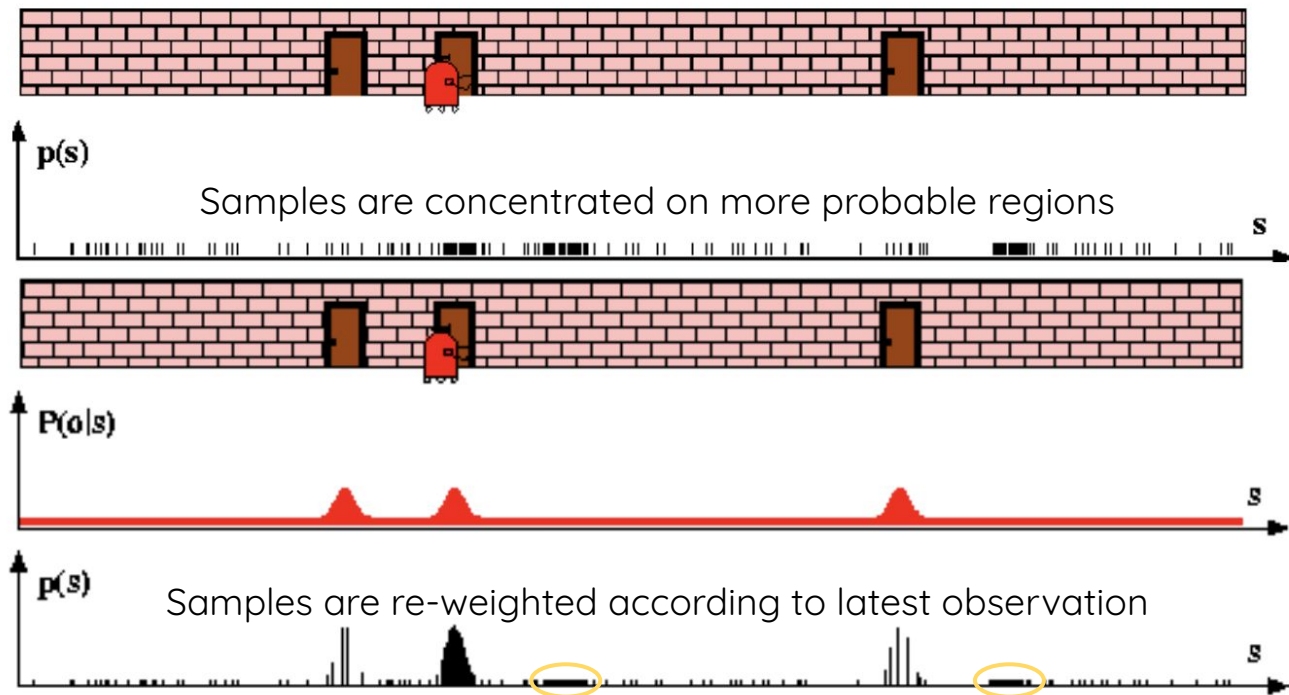
- Instead of sampling from target distribution \mathbf{f} directly, sample from a proposal distribution \mathbf{g}
- Then, the samples are weighted according to \mathbf{f}

$$w^{[m]} = \frac{f(x^{[m]})}{g(x^{[m]})}$$

- Selection of \mathbf{g} is important, closer to \mathbf{f} the better
- But, even for non-ideal \mathbf{g} , the final weights resemble \mathbf{f}
- There is one problem: Improbable regions may have the majority of the particles!
(it gets worse and worse after many iterations)




Multimodal Belief Example Revisited



Yellow regions contain a lot of samples even though not very probable, what to do?

Resampling

- The problem is also called “particle degeneracy”: Lack of particles corresponding to the true state (since most particles are in improbable regions, i.e. their weights are close to zero)
- The solution: **Resampling!**



- Resample particles according to their weights! (“survival of the fittest”)
- A particle with a high weight has a higher chance to be selected (even duplicated)
- Particles with lower weights won’t be selected at all (we’ll get rid of improbable particles)
- There are many resampling techniques: Low Variance Sampling (above example), Stochastic Universal Sampling, Roulette Wheel... (we won’t go into details)

Particle Filter Algorithm

1: **Algorithm Particle filter**($\mathcal{X}_{t-1}, u_t, z_t$):

2: $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$

3: for $m = 1$ to M do

4: sample $x_t^{[m]} \sim p(x_t \mid u_t, x_{t-1}^{[m]})$

5: $w_t^{[m]} = p(z_t \mid x_t^{[m]})$

6: $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$

7: endfor

8: for $m = 1$ to M do

9: draw i with probability $\propto w_t^{[i]}$

10: add $x_t^{[i]}$ to \mathcal{X}_t

11: endfor

12: return \mathcal{X}_t

Prediction: Advance particle to its next a priori state

Correction: Compute particle weight according to measurement

Resampling

(we don't have to do it at each iteration, in which case the weights are reset to $1/M$ here)

When to resample?

- We may do it at every step, but what if we lose an important particle by chance?
- Resampling is a destructive operation, improves our estimate but alters it
- Resampling is also a computationally heavy operation
- A better strategy is to do it based on **effective sample size**:

$$\hat{N}_{\text{eff}} = \frac{1}{\sum_{i=1}^N (\tilde{w}_k^{(i)})^2}$$

- All weights equal to each other: $N_{\text{eff}} = N$
- All weights zero except one particle with weight 1: $N_{\text{eff}} = 1$
- Idea: Resample only when effective sample size is lower than a threshold

Particle Depletion

- Even if we do resampling, we may still end up with no particles around the true state:
Maybe due an unlucky streak of bad measurements...
- The solution: **Particle Injection**
- Idea: Add a small amount of additional particles randomly
 - Could be generated around the latest measurement
 - Could be generated totally randomly (similar to prior belief)
- Advantage is its simplicity, but theoretically we are altering our posterior belief
(but negligible if injected particles are few)
- A famous problem: Kidnapped Robot Problem
 - What if a robot is carried to a different location?
 - We won't be able to trust the motion model (since robot is carried)
 - We want to be able to correct our estimate towards the true state: particle injection is a solution

Practical Applications of Data Fusion


- We will now explore some applications of fusion methods for positioning
- We've already given some examples that can be readily applied to the real world (1D, 2D positioning...), but there are some additional considerations
- How to convert inertial sensor measurements to viable inputs?
- How to cope with inertial sensor errors?

Using Inertial Sensors for the Motion Model

- Inertial navigation equations can be used as is for gyroscopes
 - Errors are manageable, only one integration required
- We can't say the same thing for accelerometer...
 - If you have frequent and reliable correction measurements, you can still do
Example: RTK (Real-Time Kinematics) systems -> High quality, high frequency GPS + inertial navigation equations (double integration...)
 - Frequent corrections mean small delta-time, which means double integration error remains bounded
 - Otherwise, not feasible
 - But there are some options for **pedestrian dead reckoning**...



Pedestrian Dead Reckoning

- Assuming we are tracking walking humans
- Human gait has some properties that we can exploit



Zero-Velocity Updates (ZUP)

- Applicable for sensors fixed to the foot (e.g. sensors embedded in shoes, or simply IMUs strapped to the feet)
- There is a brief phase where foot is stationary within the gait cycle
 - Foot is completely flat on the floor, velocity is known to be zero
- Detectable via raw gyro readings via thresholding (also possible with accelerometer)



Error-State Kalman Filter

- Rather than tracking the actual state via KF, track the error
- Actual state (position, velocity, attitude) is tracked via standard navigation equations
- ZUP measurements are used to correct velocity
- KF automatically corrects position as well!
(e.g. if velocity overshoot, position will be corrected backwards)

PREDICT

$$\mathbf{P}_{t|t-1} = \mathbf{F}_t \mathbf{P}_{t-1|t-1} \mathbf{F}_t^\top + \mathbf{\Sigma}_w$$

CORRECT





$$\mathbf{K}_t = \mathbf{P}_{t|t-1} \mathbf{H}^\top \left(\mathbf{H} \mathbf{P}_{t|t-1} \mathbf{H}^\top + \mathbf{\Sigma}_v \right)^{-1}$$

$$\delta \hat{\mathbf{x}}_t = \mathbf{K}_t (\mathbf{0} - \mathbf{H} \hat{\mathbf{x}}_t)$$

$$\mathbf{P}_{t|t} = \mathbf{P}_{t|t-1} - \mathbf{K}_t \mathbf{H} \mathbf{P}_{t|t-1}$$


- ZUP is marked with red: just a vector of zeros!
(in 3 axes)
- NOTE: No need to track mean (\mathbf{x}), since mean error is always zero, all we need is covariance
- $\delta \hat{\mathbf{x}}_t$ is then applied to the separately tracked actual state

Zero-Velocity Updates Example




Inertial Tracking for Mobile Phones

- ZUP is not possible, since phones are usually not strapped to foot...
- Assuming the phone remains more or less stationary with respect the the user (e.g. handheld, strapped to a belt, in the pocket)
- Attitude can be projected to the Earth-tangent plane (converted to heading/yaw)
- How to estimate position?
- **Step detection** to the rescue!




Step Detection



- Human gait is periodic
- Naive approach: Peak/valley detection
 - Typically performed on the magnitude of acceleration (no need for attitude estimation even!)
 - If we can compute attitude, we can extract only the z component (may be more reliable)
 - Works really well, but we've gotta be careful about false positives!
- Detecting steps is not enough, we also have to estimate step length...

Step Length Estimation

- All we can do is approximate, but hopefully typical step length remains in a limited range
 - We can assume a fixed step length
 - May be good enough provided we have accurate correction measurements
- Filters can learn the undershoot/overshoot!
 - Start with a fixed estimate, it gets better over time
 - We should include a term for this in the filter state vector (aka **bias term**)
 - The same concept can be applied to learn other sensor bias errors!
 - Bias estimation is a difficult task though
 - Correction measurements should be accurate
- Not ideal initially, can we do better?



Step Length Estimation

- There are existing approximation models that can give us a somewhat better estimate
- Intuition: Step length influences acceleration (e.g. longer steps tend to produce more acceleration)
- Two very commonly used models:

WEINBERG

$$k \cdot \sqrt[4]{a_{\max} - a_{\min}}$$


SCARLET

$$k \cdot \frac{\left(\sum_{i=1}^N |a_i| / N\right) - a_{\min}}{a_{\max} - a_{\min}}$$

- max/min acceleration are determined within the window of a detected step
- ***k*** is a tunable coefficient
- Time could also be incorporated into the model (e.g. longer steps tend to be quicker)
- Not foolproof, but may be better than a fixed estimate
- Could be similarly augmented via bias estimation (for instance we could estimate ***k***)


Particle Filter Positioning Example

- A simple motion model consisting of displacement + heading change
 - Similar to velocity + heading change model, could easily be converted to/from
 - Suitable for step detector output (step length = displacement)
- Particles are propagated with these relative measurements (+ some noise)




Particle Filter: How to incorporate map info?

- Maps have unreachable regions such as walls, can we use this to improve our estimate?
- A highly non-linear problem, posterior distribution can become very complex
- Solution: Kill particles that bump into walls (or make their weights low)




Particle Filter: Correction Example

- An absolute position measurement can then be used for correction
- Particle weights are updated according to the proximity to the new measurement
- Measurement model could be Gaussian, or anything really



Particle Filter: Resampling Example


- Particles are then resampled; the most probable particles are duplicated
- Weights are reset to be equal
- Initially the duplicated particles are in the same position, but in the next predict iteration, random noise will separate them



Real World Data Fusion Example

- Tracking a hand-held phone
- Inertial data: Step detector + gyro-based heading estimation (no magneto)
- Absolute position measurements: Bluetooth beacons trilateration
- Particle filter for fusion (state: x, y, heading)
- Also uses map information (obstacles) to improve estimate
- We'll see **inertial only**, **Bluetooth beacons only**, and finally **fusion**


Inertial Only



Bluetooth Beacons Only



Fusion



Fusion

Absolute

- Good for long-term (> 1 minute)
- Anchored on the map

Beacon-Based


Relative

- Good for short-term (seconds)
- Not anchored on the map

Inertial

FUSION

Location



Detecting Floor

- Indoor positioning demands floor detection, as buildings have multiple floors
 - A single coordinate may correspond to multiple floors, which is it?
- Radio-based positioning can give an idea
 - Lateration/angulation can operate on 3D (if you are using a suitable tech)
 - The floor info can be embedded into the fingerprinting map
 - The receiver tends to pick up beacons/AP on the same floor more strongly
- Indoor spaces can be very challenging: Mezzanines, vertical spaces, inclined surfaces...
 - Radio-based methods may be insufficient
 - Fluctuations may make it difficult to converge on a solution (especially for RSSI-based methods)
- An augmenting solution: **Barometer!**

Barometer

- Most phones have a digital barometer nowadays
- Measures atmospheric pressure
 - Measurements are absolute, however pressure at the same location/altitude changes according to weather, time of day...
- Surprisingly accurate for measuring relative altitude changes
 - We can roughly convert the difference between two pressure values to relative altitude change
 - Accurate on the order of < 1 meter
- Relative altitude change can be incorporated into a motion model!
 - Can then be fused with other radio-based data
 - Can handle momentary fluctuations and bias
 - Should be careful on how much to trust it (what if our first estimate is wrong?)



Conclusion

- Data fusion methods can be used to combine complementary data (one filling in the weaknesses of the other)
- Fusion can take place either in the motion model (as control input), or in the measurement model (multiple measurement models are possible!)
- Non-linearity and the posterior distribution characteristics are important for method selection
- Fusion methods are only as powerful as how accurate you select your error models!
- Particle filters are very powerful, but one should be mindful of the computation cost; for many problems a linearized KF may be just as good!